

# Mercari's Large Scale System Design Hackathon

**Team 5**



Inter-IIT Tech  
Meet 10.0



**mercari**

# Design Documentation

## 1. Requirement Specifications

- 1.1. Patient Registration
- 1.2. Hospital, Lab and Pharmacy Inventory Management
- 1.3. Doctor Appointment Scheduling.

## 2. Design Decisions

### 2.1. Assumptions and Trade-offs

- We are using Kubernetes for container orchestration, and docker as the underlying containerization software.
- We are making sure that the entire hospital chain is composed of decoupled microservices, which allows us to scale the services independently.
  - One of the drawbacks of the microservices approach is the service calls taken to serve a request have substantial latency.
- We are using normalized relational databases, for maintaining consistency and removing redundancy.
- We are using a centralized auth service, which makes sure all of the service calls are authenticated and authorized.
- The entire Kubernetes cluster is monitored for all the relevant metrics, making it easy for the developers to identify and track the issues with the service calls if needed.
- We are using Sentry for managing centralized error logging, making managing and tracking errors much simpler, which helps to audit our application with ease.
- We are using AWS RDS as the centralized database and AWS EKS for the Kubernetes Cluster. Using a managed service helped in cutting short of the development time while providing scalability and high availability out of the box.
- We chose AWS over GCP, due to the fine-grained control AWS provides. GCP has an opinionated approach, which helps to simplify the development, but we need more custom control in the Kubernetes deployment.

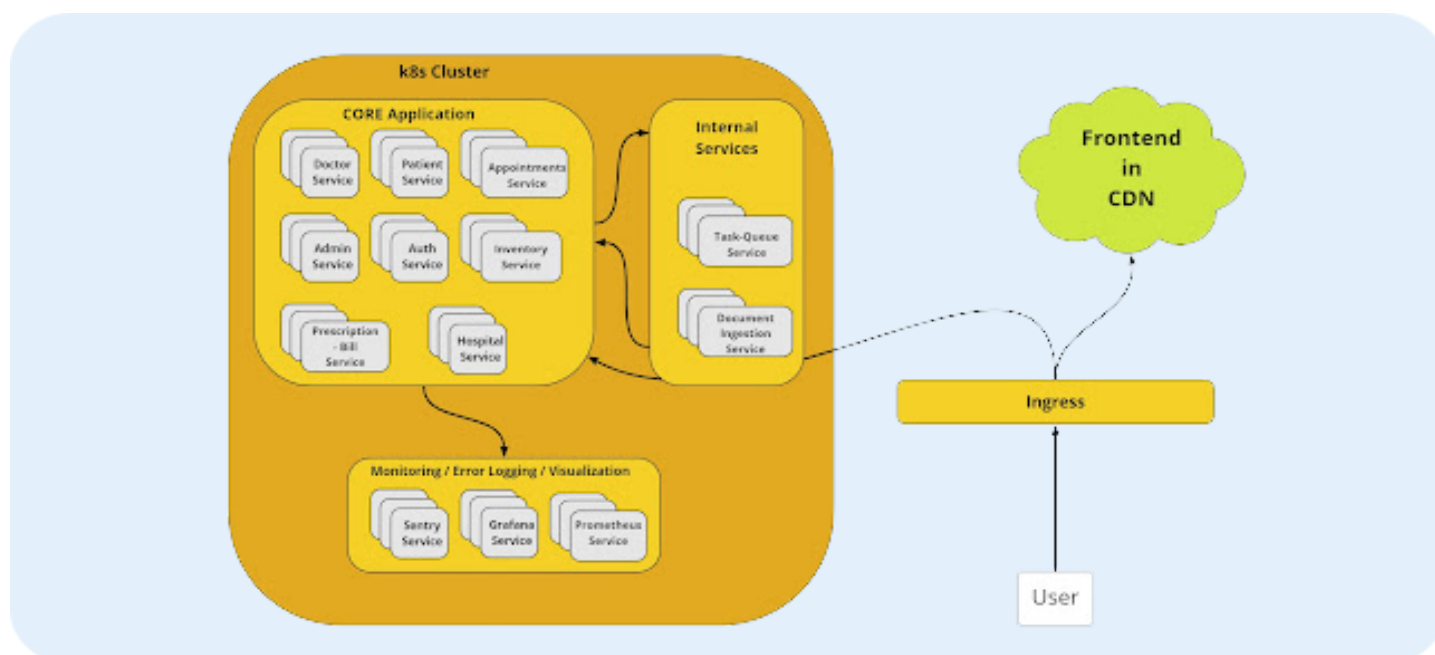
### 2.2. Tech & Infrastructure Stack

For the services implemented, we are using -

- ◆ NodeJS, ExpressJS for hospital and patient service.
- ◆ Python, Flask-RestX, and uWSGI/Gunicorn for the rest of the micro-services.
- ◆ MySQL as the RDBMS, and Redis as an in-memory database.

- ◆ Celery with a distributed RabbitMQ deployment as a message broker for a task queue to carry out asynchronous tasks.
- For the Infra side of things, we are using AWS EKS, Prometheus, Sentry, Grafana, AWS S3, AWS RDS.
- The S3 ingestor service is public cloud-agnostic, we can plugin any object storage as we wish. We have added the S3 plugin as one of the cloud services.

## 2.3. System Architecture Overview



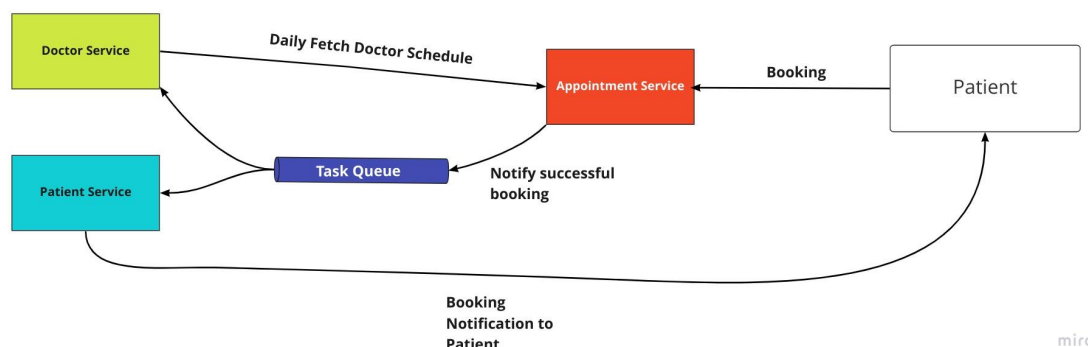
- We have broken down the entire project cluster into multiple microservices, which helps us in decoupling individual services and scaling them independently based on our requirements
- We have an Open Source **Integrated Monitoring/Error Logging/Visualization Software Suite**, which helps us in *identifying/tracking down/visualization* the metrics associated with the project
  - Having Open Source software dependency helps us in preventing vendor-lock-in, and ability to get latest updates/feature additions for the software.
- The entire setup is based on Kubernetes, so we can easily scale horizontally.
- We have custom dashboards for visualizing the cluster, and the metrics associated. This helps us in identifying the health of the cluster, and provision more Nodes if needed.

## 2.4. Important Flows

### → Auth Flow

- ◆ Given the requirement of compatibility with mobile apps, JWT is the best authentication mechanism to use.

- ◆ Using JWT helps us to get rid of Load Balancer affinities that affect traditional Cookie-based systems that rely on the login data stored in the RAM of the server.
- ◆ We use a global secret that stores the JWT Secret key. Hence, the secret key is shared across multiple microservices.
- ◆ A separate Auth service handles login, registration, and refreshing of tokens. It also handles the blacklisting of tokens and provides APIs for checking this blacklist through inter-service API calls.



## → Appointment Saga

- ◆ We notice that the number of users catered by the doctor service and the appointment service vary by an order of magnitude of 50-100, hence, they need to be scaled separately. For example, in bigger cities, we have a lot of patients that come during weekends.
- ◆ The appointment service is configured to pull the schedule of each doctor on a daily basis and then book patients in the time slot based on this pulled view of the schedule.
- ◆ As patients book slots, the booking event is pushed to a Task Queue that asynchronously messages the doctors and patient regarding the confirmation of booking.
- ◆ This allows us to work with an **eventually consistent** appointment scheduling system, which is perfect for our use case.

## → Storage Access Control

- ◆ We implement Role-Based Access Control (RBAC) for the prescriptions, lab reports, and medical history files uploaded by the users of the system.
- ◆ We implement this by creating a thin wrapper over S3 which checks the ownership of files using metadata stored in a Redis database service.
- ◆ This storage wrapper can be easily extended to any other object storage system as well.

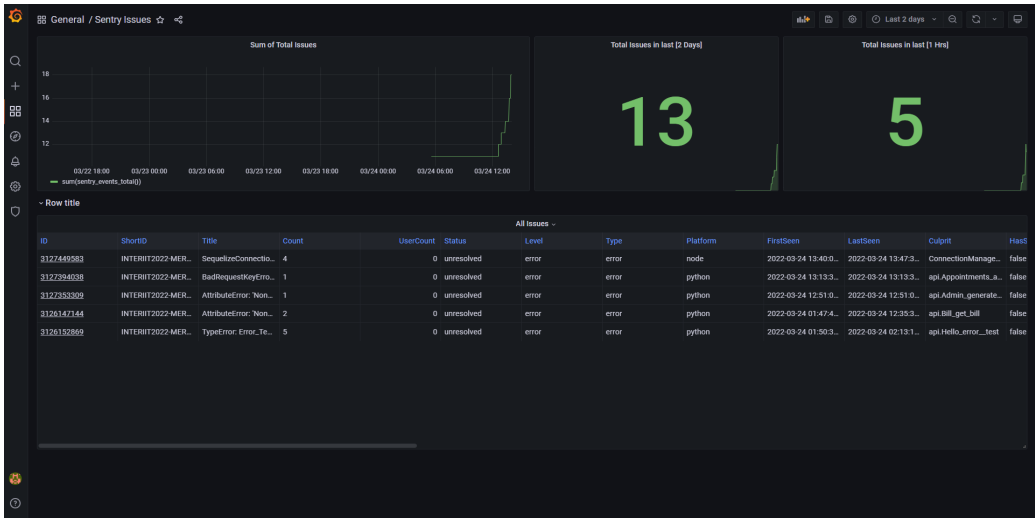
## → Monitoring, Error Logging and Visualization

- ◆ *We understand that monitoring and visualization are an important part of running large-scale distributed systems. Having a transparent overview of the health of the application is really important.*
- ◆ We are using the Grafana Dashboard for rich visualizations of the cluster metrics.
- ◆ We have Service Level Metrics for showing the 2xx/3xx/5xx errors for each of the requests handled.
- ◆ We have also added p95/p99 latency for getting an idea of the latency for each service request.

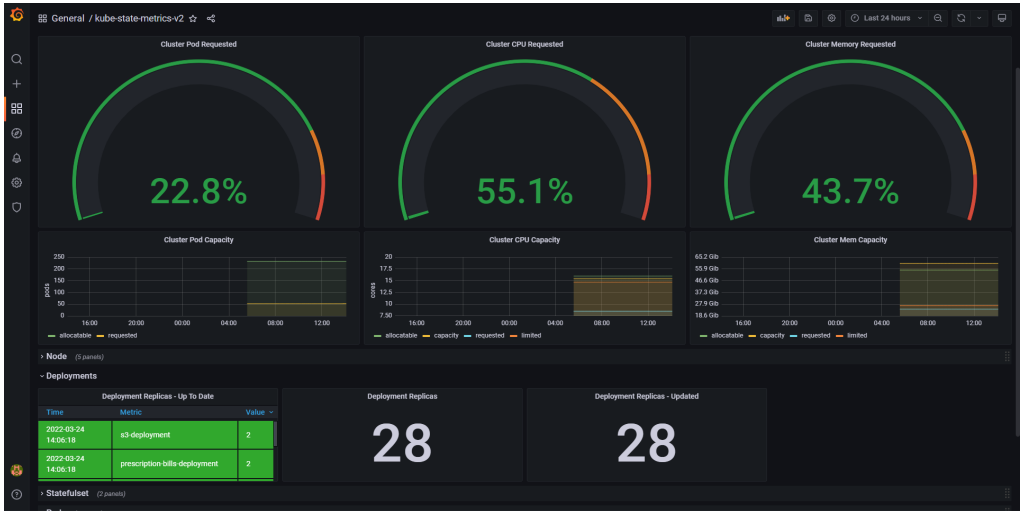
- ◆ We have another dashboard for getting relevant metrics for the cluster, which helps us in provisioning more Nodes, in case the requested percentage(%) of resources crosses a threshold.
- ◆ We have an uptime dashboard for getting the percentage(%) of available replicas for the service.



Service Level Metrics



Centralized Error Logging Dashboard



Kube-state-metrics: Kubernetes Cluster Metrics

### 3. Scalability and Portability

- The use of Kubernetes allows us to cater to the national population easily, scaling up to as much as India's population, while maintaining a 99.99% of availability. However, for a growing user base, our modular design allows for easy swapping of components.
- In particular, we can replace the managed MySQL instances we use for our database with more resilient CockroachDB for geo-distributed scaling.
- We have taken special care to build our solution solely based on Open Source components with very minimal vendor lock-in. Our solution can be deployed in any popular Kubernetes service (EKS, GKE, AKS) or on-premises through OpenStack.

### 4. Cost Analysis

We estimate that for a national level usage, we would need about 20 m5.large EKS nodes with 5 db.t4g.large MySQL RDS instances and 100GB/day of data egress and about 50TB data stored in s3. These are the major players in the costing landscape.

*Since the budget is of the utmost priority while designing a system for government hospital chains, we have tried to make this as cost-efficient as possible.*

Given these constraints, the daily cost of running the application on a national scale is as follows (as per current AWS rates in ap-south-1):

Resource	Rate	Quantity	Total price per day
m5.large instances	\$0.101/hr	20	\$48.48
db.t4g.large	\$0.167/hr	5	\$20.04
Egress	\$0.1093/GB	100 GB/day	\$10.93
NAT Gateway	\$0.056/GB	100 GB/day	\$5.60
S3	\$0.025/GB	50 TB/month	\$42.67
EKS Cluster charges	\$0.1/hr	1	\$2.4
		<b>Total</b>	\$130.12 (~ ₹ 10,000)

### 5. Deliverables

We have exposed all our services to the external network for evaluation purposes. The links for accessing the Swagger UI for API testing are as follows :

admin	<a href="http://a9c092961e2974c9c9aa6d87d6ec23ad-1318982453.ap-south-1.elb.amazonaws.com">http://a9c092961e2974c9c9aa6d87d6ec23ad-1318982453.ap-south-1.elb.amazonaws.com</a>
appointments	<a href="http://a2b9b9bef1490fa48d18825d04f399814b-1958880394.ap-south-1.elb.amazonaws.com">http://a2b9b9bef1490fa48d18825d04f399814b-1958880394.ap-south-1.elb.amazonaws.com</a>



auth	<a href="http://aof4b292f577243a5971d04ff7da5b94-987059869.ap-south-1.elb.amazonaws.com">http://aof4b292f577243a5971d04ff7da5b94-987059869.ap-south-1.elb.amazonaws.com</a>
doctor	<a href="http://a3d4df110319947cda7c3e950194e806-1836406126.ap-south-1.elb.amazonaws.com">http://a3d4df110319947cda7c3e950194e806-1836406126.ap-south-1.elb.amazonaws.com</a>
hospital-crud	<a href="http://aaf792c0a22124ee59a68c5a90407979-1714793882.ap-south-1.elb.amazonaws.com">http://aaf792c0a22124ee59a68c5a90407979-1714793882.ap-south-1.elb.amazonaws.com</a>
inventory	<a href="http://aee59aa5b31254a1785b8e39ddc1814a-1896607215.ap-south-1.elb.amazonaws.com">http://aee59aa5b31254a1785b8e39ddc1814a-1896607215.ap-south-1.elb.amazonaws.com</a>
patient-crud	<a href="http://a5503b4300e3742af8d8b9197298b705-2124458015.ap-south-1.elb.amazonaws.com">http://a5503b4300e3742af8d8b9197298b705-2124458015.ap-south-1.elb.amazonaws.com</a>
prescription-bills	<a href="http://af0fdf42db6fb4bcc9b460c155384d86-289534987.ap-south-1.elb.amazonaws.com">http://af0fdf42db6fb4bcc9b460c155384d86-289534987.ap-south-1.elb.amazonaws.com</a>
phpmyadmin	<a href="http://a3b2ed2f68be0488f90562aec2f95970-54686662.ap-south-1.elb.amazonaws.com">http://a3b2ed2f68be0488f90562aec2f95970-54686662.ap-south-1.elb.amazonaws.com</a>
ingestion (s3)	<a href="http://ae18803d8fef4493b984b56bb3c8b050-579456122.ap-south-1.elb.amazonaws.com">http://ae18803d8fef4493b984b56bb3c8b050-579456122.ap-south-1.elb.amazonaws.com</a>
task-queue (tq)	<a href="http://a4daf9353217547d3a283965ab13367a-121713899.ap-south-1.elb.amazonaws.com">http://a4daf9353217547d3a283965ab13367a-121713899.ap-south-1.elb.amazonaws.com</a>
grafana	<a href="http://a35533eeebc134e2a8fd73739b48a532-1813051043.ap-south-1.elb.amazonaws.com">http://a35533eeebc134e2a8fd73739b48a532-1813051043.ap-south-1.elb.amazonaws.com</a>
prometheus	<a href="http://a238d1c38a3d14b54bdb92481d53a12f-1330506167.ap-south-1.elb.amazonaws.com/">http://a238d1c38a3d14b54bdb92481d53a12f-1330506167.ap-south-1.elb.amazonaws.com/</a>

**Grafana Username:** admin

**Grafana Password:** QKrcgxbGrTliNLTv55zcZqrUOXv4BzGg1CJXJAZa

**Frontend URL:** <http://mercari-interiit-website.s3-website.ap-south-1.amazonaws.com/>

## 6. Future Enhancements

- We can integrate tracing for the entire Kubernetes setup -
  - Tracing provides us the full path of the service calls, and with inter-micro-services communication, it becomes difficult to debug and reproduce the errors.
  - Having a tracing infrastructure helps us in quickly identifying issues
- Having a centralized logging setup, like grafana Loki, would help in accessing/aggregating the logs for all of the micro-services in one place.
- Configuring Prometheus Alerts, in case of a node crash, unexpected rise in CPU/Memory Usage for every deployment/node/pod, unexpected increase in latency, unavailability of microservices.
- We can set up Error Budgets / Service Level Objectives / Service Level Agreements for each of the services, to make sure we have transparency with the clients about the real-time performance of the services.

## 7. References

1. <https://microservices.io/patterns/data/saga.html>
2. <https://docs.celeryq.dev/en/stable/>
3. <https://prometheus.io/>
4. <https://grafana.com/>
5. <https://sentry.io/welcome/>
6. <https://aws.amazon.com/eks/>
7. <https://redis.io/>
8. <https://www.cockroachlabs.com/>